

Airbnb - Boston

```
In [1]: # Import Neccessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import math
import time

# Import library for VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.linear_model import ElasticNet
from sklearn import preprocessing
from sklearn.metrics import *

import folium
from folium import plugins
from folium.plugins import HeatMap
```

```
In [2]: ur = "http://data.insideairbnb.com/united-states/ma/boston/2021-12-17/data/"
url = "http://data.insideairbnb.com/united-states/ma/boston/2021-12-17/visu
```

We have imported pandas ,numpy for basic data analysis. Seaborn and matplotlib for data visualization. Folium for maps.

```
In [3]: # Loading Dataset
boston = pd.read_csv(ur,compression='gzip',low_memory=False)
boston2 = pd.read_csv(url)
```

Data Exploration , Visualization and Processing

```
In [4]: ### Created a Subset of data
dat = boston[['host_id', 'accommodates', 'bedrooms', 'beds', 'amenities']].copy
boston3 = pd.merge(boston2, dat)
boston3.tail()
```

Out[4]:

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitu
107384	53772114	Entire Modern Apartment in Downtown Boston 2 R...	434990435	Luke	NaN	South End	42.3448
107385	53772114	Entire Modern Apartment in Downtown Boston 2 R...	434990435	Luke	NaN	South End	42.3448
107386	53772114	Entire Modern Apartment in Downtown Boston 2 R...	434990435	Luke	NaN	South End	42.3448
107387	53772114	Entire Modern Apartment in Downtown Boston 2 R...	434990435	Luke	NaN	South End	42.3448
107388	53756737	The Arbor Retreat	42715907	Emma	NaN	Jamaica Plain	42.3048

5 rows × 22 columns



```
In [5]: ### Calculating Summary Statistics
boston3.describe()
```

Out[5]:

	id	host_id	neighbourhood_group	latitude	longitude	pr
count	1.073890e+05	1.073890e+05	0.0	107389.000000	107389.000000	107389.000000
mean	4.461479e+07	1.758024e+08	NaN	42.342998	-71.084340	147.847
std	1.192542e+07	1.221205e+08	NaN	0.019351	0.033584	116.540
min	3.781000e+03	4.804000e+03	NaN	42.235330	-71.172520	0.000
25%	4.208014e+07	1.074344e+08	NaN	42.335150	-71.101850	50.000
50%	4.961078e+07	1.074344e+08	NaN	42.347970	-71.070840	149.000
75%	5.227869e+07	2.978601e+08	NaN	42.355180	-71.061640	212.000
max	5.383997e+07	4.349904e+08	NaN	42.392790	-70.997810	10000.000

```
In [6]: ### Precise Summary of the Dataframe
boston3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107389 entries, 0 to 107388
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    107389 non-null  int64
1   name                                107389 non-null  object
2   host_id                             107389 non-null  int64
3   host_name                           61820 non-null   object
4   neighbourhood_group                 0 non-null       float64
5   neighbourhood                       107389 non-null  object
6   latitude                           107389 non-null  float64
7   longitude                           107389 non-null  float64
8   room_type                           107389 non-null  object
9   price                               107389 non-null  int64
10  minimum_nights                      107389 non-null  int64
11  number_of_reviews                   107389 non-null  int64
12  last_review                         29052 non-null   object
13  reviews_per_month                   29052 non-null   float64
14  calculated_host_listings_count      107389 non-null  int64
15  availability_365                     107389 non-null  int64
16  number_of_reviews_ltm               107389 non-null  int64
17  license                             23450 non-null   object
18  accommodates                        107389 non-null  int64
19  bedrooms                            88219 non-null   float64
20  beds                                82280 non-null   float64
21  amenities                           107389 non-null  object
dtypes: float64(6), int64(9), object(7)
memory usage: 18.8+ MB
```

```
In [7]: ### Check for null values  
boston3.isna().sum()
```

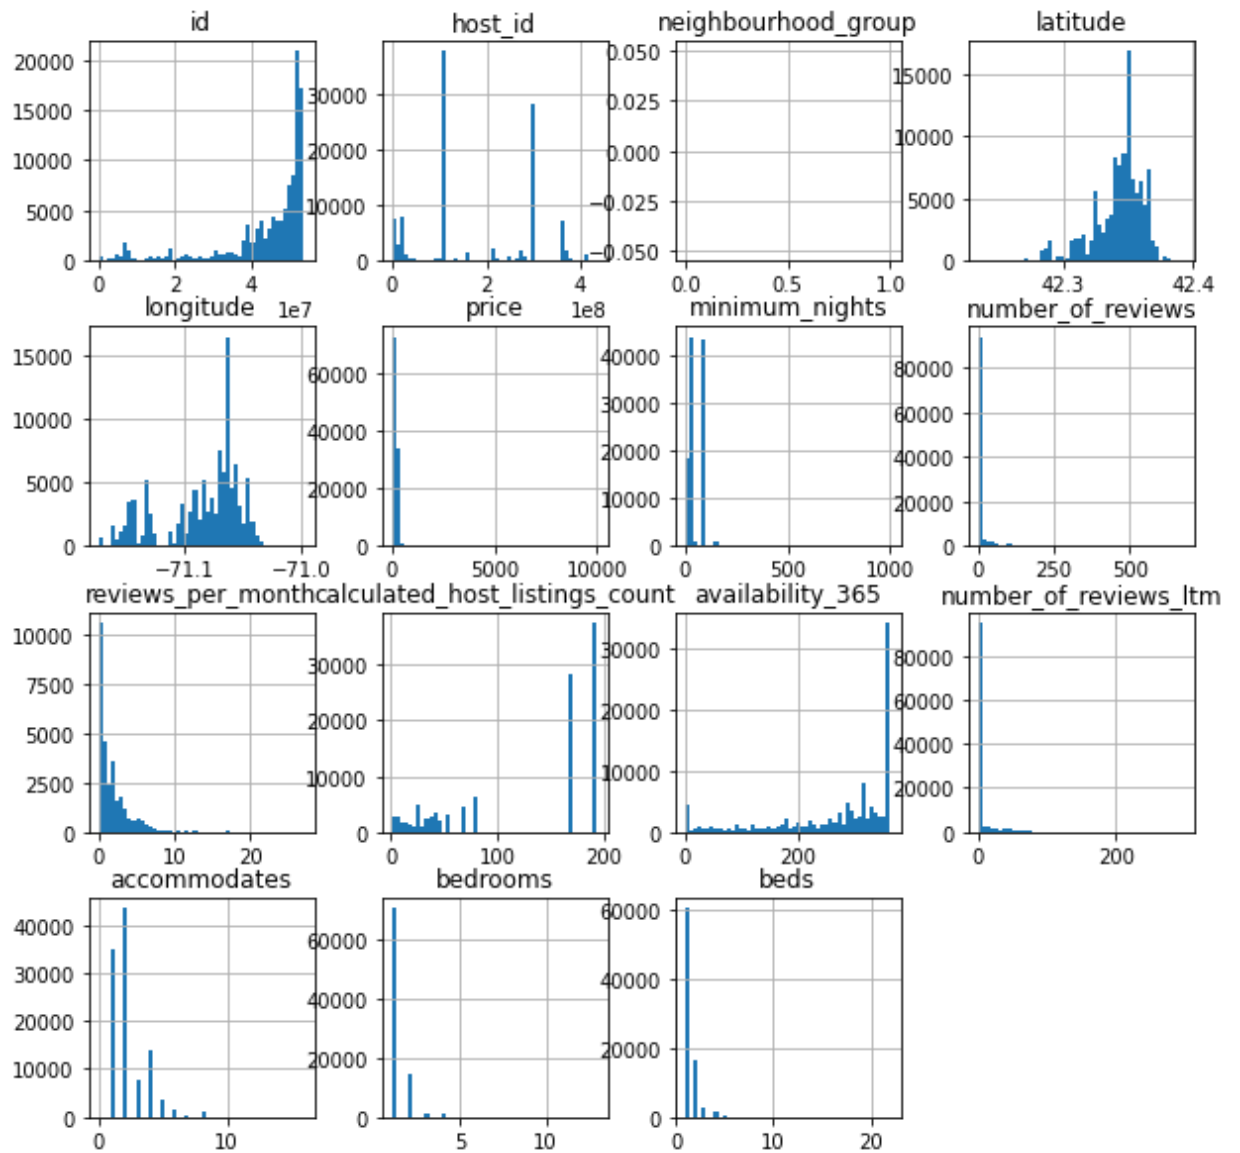
```
Out[7]: id          0  
name            0  
host_id        0  
host_name      45569  
neighbourhood_group  107389  
neighbourhood    0  
latitude        0  
longitude       0  
room_type       0  
price          0  
minimum_nights  0  
number_of_reviews  0  
last_review     78337  
reviews_per_month  78337  
calculated_host_listings_count  0  
availability_365  0  
number_of_reviews_ltm  0  
license        83939  
accommodates    0  
bedrooms       19170  
beds           25109  
amenities       0  
dtype: int64
```

```
In [8]: ### Percentage of null rows in each column
def missing(datas):
    print (round((datas.isnull().sum() * 100 / len(datas)),2).sort_values(ascending=True))

missing(boston3)
```

neighbourhood_group	100.00
license	78.16
last_review	72.95
reviews_per_month	72.95
host_name	42.43
beds	23.38
bedrooms	17.85
id	0.00
accommodates	0.00
number_of_reviews_ltm	0.00
availability_365	0.00
calculated_host_listings_count	0.00
number_of_reviews	0.00
name	0.00
minimum_nights	0.00
price	0.00
room_type	0.00
longitude	0.00
latitude	0.00
neighbourhood	0.00
host_id	0.00
amenities	0.00
dtype: float64	

```
In [9]: boston3.replace({'f': 0, 't': 1}, inplace = True)
boston3.hist(bins=50, figsize=(10,10))
plt.savefig('distribution.png', dpi=650, bbox_inches='tight')
plt.show()
```



We can drop the columns with less categories. Checking whether boolean and categorical features contain sufficient numbers of instances in each category to make them worth including. It can be seen that several columns only contain one category and can be dropped while preprocessing.

Data Preprocessing

```
In [10]: # To protect the privacy of the hosts and reviewers, we drop them as well as  
boston3.drop(['id', 'host_name', 'last_review', 'license'], axis=1, inplace=True)
```

```
In [11]: boston3.drop(['neighbourhood_group'], axis=1, inplace=True)
```

```
In [12]: boston3.isnull().sum()
```

```
Out[12]: name                0  
host_id                    0  
neighbourhood              0  
latitude                   0  
longitude                  0  
room_type                  0  
price                      0  
minimum_nights             0  
number_of_reviews          0  
reviews_per_month          78337  
calculated_host_listings_count  0  
availability_365           0  
number_of_reviews_ltm      0  
accommodates               0  
bedrooms                   19170  
beds                       25109  
amenities                  0  
dtype: int64
```

```
In [13]: #Percentage of null rows in each column
def missing(datas):
    print (round((datas.isnull().sum() * 100 / len(datas)),2).sort_values(ascending=True))

missing(boston3)
```

reviews_per_month	72.95
beds	23.38
bedrooms	17.85
name	0.00
accommodates	0.00
number_of_reviews_ltm	0.00
availability_365	0.00
calculated_host_listings_count	0.00
number_of_reviews	0.00
host_id	0.00
minimum_nights	0.00
price	0.00
room_type	0.00
longitude	0.00
latitude	0.00
neighbourhood	0.00
amenities	0.00
dtype: float64	

```
In [14]: ### Filling the null values with mean and median

boston3['reviews_per_month'].fillna(boston3['reviews_per_month'].mean(), inplace=True)
boston3['beds'].fillna(boston3['beds'].median(), inplace=True)
boston3['bedrooms'].fillna(boston3['bedrooms'].median(), inplace=True)
```





```
In [15]: missing(boston3)
```

name	0.0
reviews_per_month	0.0
beds	0.0
bedrooms	0.0
accommodates	0.0
number_of_reviews_ltm	0.0
availability_365	0.0
calculated_host_listings_count	0.0
number_of_reviews	0.0
host_id	0.0
minimum_nights	0.0
price	0.0
room_type	0.0
longitude	0.0
latitude	0.0
neighbourhood	0.0
amenities	0.0
dtype: float64	

```
In [16]: bosub = boston3[["price", "number_of_reviews", "reviews_per_month", "calculated_host_listings_count"]]
```



```
In [17]: ### Plot to check the outliers  
boston1=bosub.select_dtypes(exclude=['object'])  
  
for column in boston1:  
    plt.figure(figsize=(10,2))  
    sb.boxplot(data=boston1, x=column)
```



```
In [18]: #Since our target variable is the price, we have standardized the values su  
# are in one standarad deviation  
x_data = boston3.price  
standard = preprocessing.scale(x_data)  
print(standard)  
  
[-0.19604426 -0.41914292 -0.41914292 ... -0.07591421 -0.07591421  
 -0.30759359]
```

```
In [21]: boston3.loc[boston3['amenities'].str.contains('Air conditioning|Central air
boston3.loc[boston3['amenities'].str.contains('Gym|24-hour fitness center|P
boston3.loc[boston3['amenities'].str.contains('Apple TV|Game console|Game c
boston3.loc[boston3['amenities'].str.contains('Avanti stainless steel elect
boston3.loc[boston3['amenities'].str.contains('Electric stove|Gas stove|Ike
boston3.loc[boston3['amenities'].str.contains('BBQ grill|Fire pit|Barbecue
boston3.loc[boston3['amenities'].str.contains('balcony|Patio or balcony|Pat
boston3.loc[boston3['amenities'].str.contains('Beach view|Beachfront|Lake a
boston3.loc[boston3['amenities'].str.contains('Bed linens|Bed sheets and pi
boston3.loc[boston3['amenities'].str.contains('Breakfast|Complimentary cont
boston3.loc[boston3['amenities'].str.contains('Coffee maker|Espresso machin
boston3.loc[boston3['amenities'].str.contains('Cooking basics'), 'cooking_b
boston3.loc[boston3['amenities'].str.contains('Dryer|Dryer \\u2013 In build
boston3.loc[boston3['amenities'].str.contains('Elevator'), 'elevator']] = 1
boston3.loc[boston3['amenities'].str.contains('Children\\u2019s books and t
boston3.loc[boston3['amenities'].str.contains('Free driveway parking on pre
boston3.loc[boston3['amenities'].str.contains('Private fenced garden or bac
boston3.loc[boston3['amenities'].str.contains('Host greets you'), 'host_gre
boston3.loc[boston3['amenities'].str.contains('Private hot tub|Shared outdo
boston3.loc[boston3['amenities'].str.contains('Internet|Pocket wifi|Wifi|Po
boston3.loc[boston3['amenities'].str.contains('Long term stays allowed'), '
boston3.loc[boston3['amenities'].str.contains('Private entrance'), 'private
boston3.loc[boston3['amenities'].str.contains('Pets|pet|Cat(s)|Dog(s)|Pets
boston3.loc[boston3['amenities'].str.contains('Safe|Security system|Securit
boston3.loc[boston3['amenities'].str.contains('Self check_in'), 'self_check
boston3.loc[boston3['amenities'].str.contains('Smoking allowed'), 'smoking_
boston3.loc[boston3['amenities'].str.contains('Step-free access|Wheelchair|
boston3.loc[boston3['amenities'].str.contains('Suitable for events'), 'even
boston3.loc[boston3['amenities'].str.contains('Aveeno body soap|Bathtub|Bee
boston3.loc[boston3['amenities'].str.contains('Sonos Bluetooth sound system
boston3.loc[boston3['amenities'].str.contains('Dedicated workspace|Dedicate
boston3.loc[boston3['amenities'].str.contains('Freezer|GE refrigerator| LG
boston3.loc[boston3['amenities'].str.contains('Onsite restaurant \\u2014 Co
```

<ipython-input-21-90ad71994ea5>:23: UserWarning: This pattern is interpreted as a regular expression, and has match groups. To actually get the groups, use str.extract.

```
boston3.loc[boston3['amenities'].str.contains('Pets|pet|Cat(s)|Dog(s)|P
ets allowed'), 'pets_allowed'] = 1
```

```
In [22]: replacecols = boston3.iloc[:,0:].columns
boston3[replacecols] = boston3[replacecols].fillna(0)
nonessential_amenities = []
for col in boston3.iloc[:,17:].columns:
    if boston3[col].sum() < len(boston3)/10:
        nonessential_amenities.append(col)
boston3.drop(nonessential_amenities, axis=1, inplace=True)
boston3.drop('amenities', axis=1, inplace=True)
```

```
In [23]: boston3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107389 entries, 0 to 107388
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   name                                  107389 non-null object
1   host_id                              107389 non-null int64
2   neighbourhood                        107389 non-null object
3   latitude                             107389 non-null float64
4   longitude                             107389 non-null float64
5   room_type                            107389 non-null object
6   price                                107389 non-null int64
7   minimum_nights                       107389 non-null int64
8   number_of_reviews                    107389 non-null int64
9   reviews_per_month                    107389 non-null float64
10  calculated_host_listings_count       107389 non-null int64
11  availability_365                      107389 non-null int64
12  number_of_reviews_ltm                 107389 non-null int64
13  accommodates                          107389 non-null int64
14  bedrooms                              107389 non-null float64
15  beds                                  107389 non-null float64
16  air_conditioning                      107389 non-null float64
17  gym                                    107389 non-null float64
18  tv                                     107389 non-null float64
19  bbq                                    107389 non-null float64
20  nature_and_views                      107389 non-null float64
21  bed_linen                             107389 non-null float64
22  coffee_machine                        107389 non-null float64
23  cooking_basics                        107389 non-null float64
24  white_goods                           107389 non-null float64
25  elevator                              107389 non-null float64
26  parking                               107389 non-null float64
27  hot_tub_sauna_or_pool                 107389 non-null float64
28  internet                              107389 non-null float64
29  long_term_stays                       107389 non-null float64
30  private_entrance                      107389 non-null float64
31  toiletries                            107389 non-null float64
32  workspace                             107389 non-null float64
33  refrigerator                           107389 non-null float64
dtypes: float64(23), int64(8), object(3)
memory usage: 28.7+ MB
```


Data Visualization

```
In [24]: boston3=boston3.drop(['nature_and_views'],axis=1)
```

```
In [25]: price_corr = boston3.corr()['price'].sort_values(ascending = False).head(6)
price_corr
```

```
Out[25]: private_entrance    0.457659
coffee_machine    0.415872
accommodates    0.413670
bed_linen    0.406062
gym    0.403006
Name: price, dtype: float64
```

```
In [26]: ### Features that influence Price
price_corr.plot(kind = 'bar', color = 'g');
plt.xticks(rotation=45);
```



```
In [27]: boston3.head()
```

```
Out[27]:
```


	name	host_id	neighbourhood	latitude	longitude	room_type	price	minimum_nights
0	HARBORSIDE-Walk to subway	4804	East Boston	42.36413	-71.02991	Entire home/apt	125	32
1	** Private! Minutes to center!**	8229	Roxbury	42.32844	-71.09581	Entire home/apt	99	3
2	** Private! Minutes to center!**	8229	Roxbury	42.32844	-71.09581	Entire home/apt	99	3
3	** Private! Minutes to center!**	8229	Roxbury	42.32844	-71.09581	Entire home/apt	99	3
4	** Private! Minutes to center!**	8229	Roxbury	42.32844	-71.09581	Entire home/apt	99	3

5 rows × 33 columns

```
In [28]: ### Barplot to find the price of different neighbourhoods in Boston

plt.figure(figsize=(15,10))
plt.xticks(rotation=50)
listings = boston3.sort_values(by = 'price')
sb.barplot(x='neighbourhood', y= 'price',palette= "bright", data = boston3)
plt.xlabel(xlabel='Neighbourhood', fontsize=15)
plt.ylabel(ylabel='price', fontsize=15)
```

Out[28]: Text(0, 0.5, 'price')



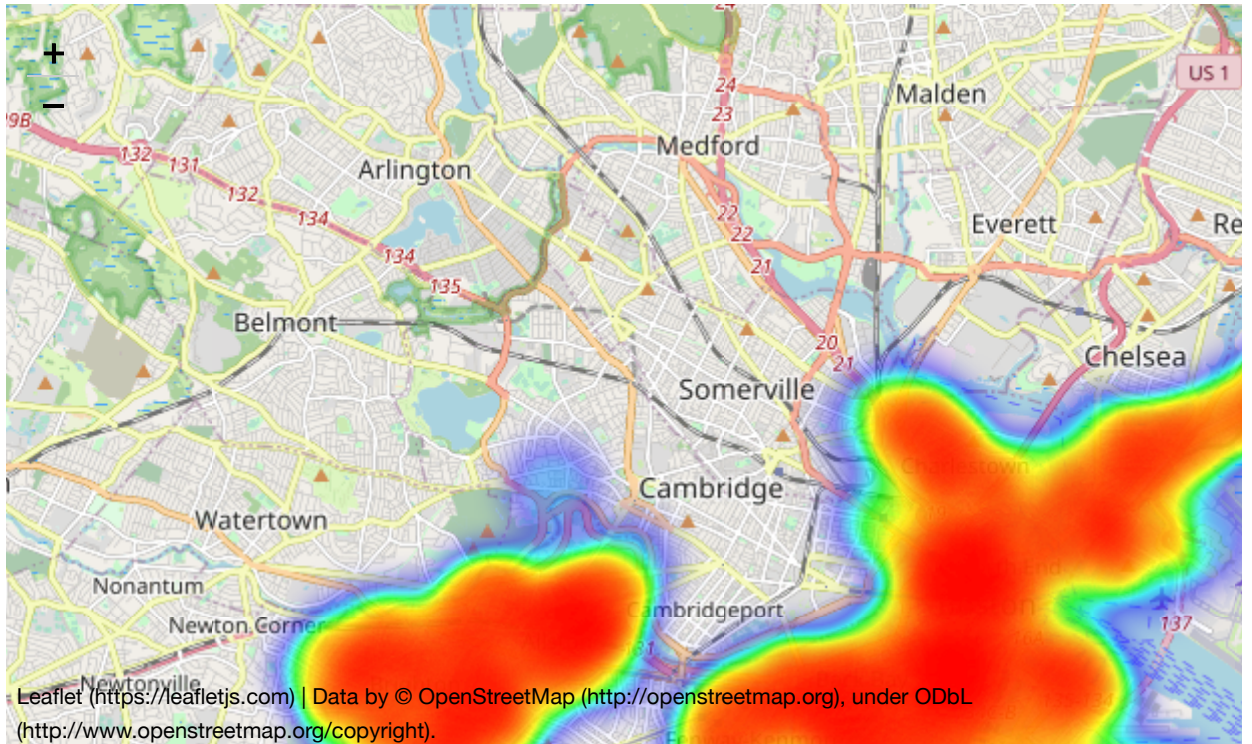
```
In [29]: maps = folium.Map(location=[42.361145, -71.057083], zoom_start = 12)

heatmap_data = [[row['latitude'],row['longitude']] for index, row in
                 boston3[['latitude', 'longitude']].iterrows()]

htm = HeatMap(heatmap_data).add_to(maps)


maps
```

Out[29]:




```
In [30]: from PIL import Image

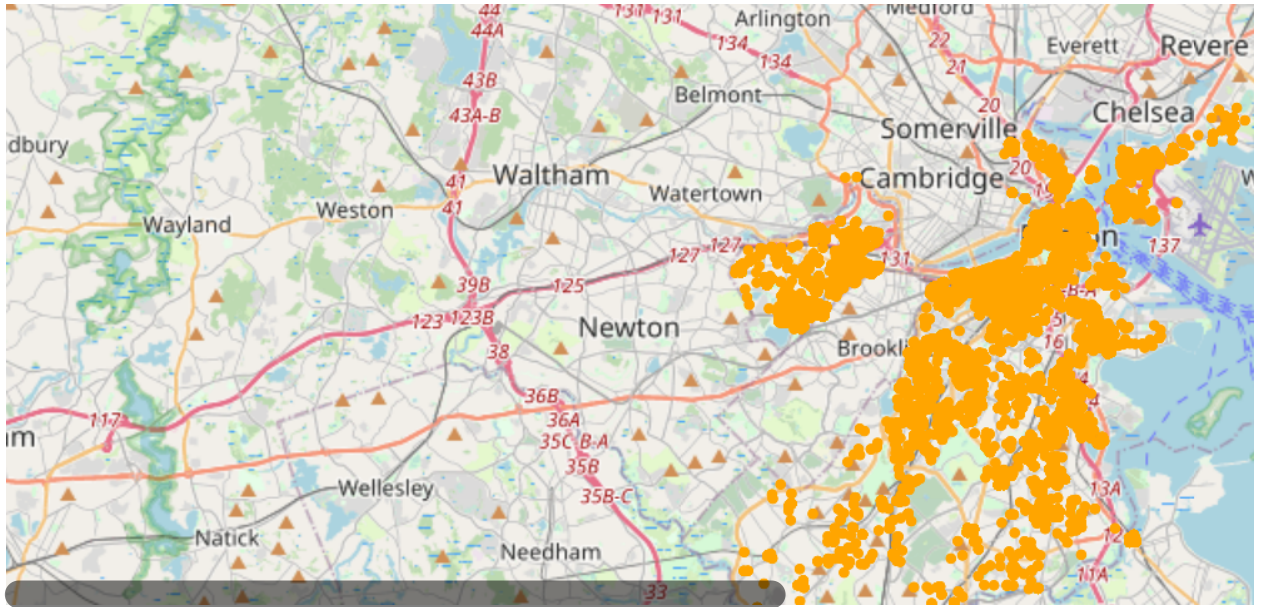
plt.figure(figsize=(13,8))
gx = plt.gca()
boston3.plot(kind='scatter',x='longitude',y='latitude',label='Listings', c=
plt.legend()
plt.show()
```




```
In [31]: ### Plot Map for Distribution of location of Airbnb properties in Boston

import plotly.express as px

fig = px.scatter_mapbox(boston3, lat="latitude", lon="longitude", hover_name="host_id",
                        color_discrete_sequence=["orange"], zoom=10, height=500)
fig.update_layout(mapbox_style="open-street-map")
fig.update_layout(margin={"r":0,"t":0,"l":0,"b":0})
fig.show()
```



```
In [32]: boston3.describe()
```

Out[32]:

	host_id	latitude	longitude	price	minimum_nights	number_of_reviews
count	1.073890e+05	107389.000000	107389.000000	107389.000000	107389.000000	107389.000000
mean	1.758024e+08	42.342998	-71.084340	147.84707	51.843345	11.1814
std	1.221205e+08	0.019351	0.033584	116.54091	36.352791	38.8621
min	4.804000e+03	42.235330	-71.172520	0.00000	1.000000	0.0000
25%	1.074344e+08	42.335150	-71.101850	50.00000	30.000000	0.0000
50%	1.074344e+08	42.347970	-71.070840	149.00000	32.000000	0.0000
75%	2.978601e+08	42.355180	-71.061640	212.00000	91.000000	1.0000
max	4.349904e+08	42.392790	-70.997810	10000.00000	1000.000000	678.0000

8 rows x 30 columns


```
In [33]: boston3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107389 entries, 0 to 107388
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   name                                  107389 non-null  object
1   host_id                              107389 non-null  int64
2   neighbourhood                        107389 non-null  object
3   latitude                             107389 non-null  float64
4   longitude                            107389 non-null  float64
5   room_type                            107389 non-null  object
6   price                                107389 non-null  int64
7   minimum_nights                      107389 non-null  int64
8   number_of_reviews                   107389 non-null  int64
9   reviews_per_month                   107389 non-null  float64
10  calculated_host_listings_count      107389 non-null  int64
11  availability_365                     107389 non-null  int64
12  number_of_reviews_ltm                107389 non-null  int64
13  accommodates                         107389 non-null  int64
14  bedrooms                             107389 non-null  float64
15  beds                                 107389 non-null  float64
16  air_conditioning                     107389 non-null  float64
17  gym                                   107389 non-null  float64
18  tv                                    107389 non-null  float64
19  bbq                                   107389 non-null  float64
20  bed_linen                            107389 non-null  float64
21  coffee_machine                       107389 non-null  float64
22  cooking_basics                       107389 non-null  float64
23  white_goods                          107389 non-null  float64
24  elevator                             107389 non-null  float64
25  parking                              107389 non-null  float64
26  hot_tub_sauna_or_pool                107389 non-null  float64
27  internet                             107389 non-null  float64
28  long_term_stays                      107389 non-null  float64
29  private_entrance                     107389 non-null  float64
30  toiletries                           107389 non-null  float64
31  workspace                            107389 non-null  float64
32  refrigerator                          107389 non-null  float64
dtypes: float64(22), int64(8), object(3)
memory usage: 27.9+ MB
```

In [34]: *### correlation verify each feature against the target feature*

```
plt.figure(figsize=(12,8))
title = 'Correlation matrix of variables'
sb.heatmap(boston3.corr(), square=True, cmap='YlGnBu')
plt.title(title)
plt.ioff()
```

Out[34]: <matplotlib.pyplot._IoffContext at 0x7f893dd399a0>



```
In [35]: boston3=boston3._get_numeric_data()
boston3=boston3.dropna(axis=0)
boston3=boston3.drop(['latitude','longitude'],axis=1)
boston3=boston3.drop(['number_of_reviews_ltm'],axis=1)
boston3.head()
```

Out [35]:

	host_id	price	minimum_nights	number_of_reviews	reviews_per_month	calculated_host_listings_c
0	4804	125	32	21	0.27	
1	8229	99	3	110	0.71	
2	8229	99	3	110	0.71	
3	8229	99	3	110	0.71	
4	8229	99	3	110	0.71	

5 rows × 27 columns

Multi-Collinearity

```
In [36]: def calc_vif(X):

    # Calculating VIF
    vif = pd.DataFrame()
    vif["variables"] = X.columns
    vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[0])]

    return(vif)
```

```
In [37]: x = boston3.iloc[:, :-1]
         calc_vif(x)
```

Out[37]:

	variables	VIF
0	host_id	7.131918
1	price	4.364104
2	minimum_nights	7.098761
3	number_of_reviews	1.909629
4	reviews_per_month	6.208965
5	calculated_host_listings_count	17.979208
6	availability_365	11.736312
7	accommodates	17.591002
8	bedrooms	30.443464
9	beds	17.887075
10	air_conditioning	6.963641
11	gym	3.731983
12	tv	11.819420
13	bbq	1.743540
14	bed_linen	9.566801
15	coffee_machine	9.450196
16	cooking_basics	4.028319
17	white_goods	10.950189
18	elevator	3.603388
19	parking	2.176904
20	hot_tub_sauna_or_pool	2.105857
21	internet	135.087303
22	long_term_stays	119.966757
23	private_entrance	6.258084
24	toiletries	1.913979
25	workspace	4.101958

```
In [38]: boston3=boston3.drop(['internet'],axis=1)
         boston3=boston3.drop(['long_term_stays'],axis=1)
         boston3=boston3.drop(['host_id'],axis=1)
```

```
In [39]: boston3.iloc[:,6:20].corr()
```

```
Out[39]:
```

	accommodates	bedrooms	beds	air_conditioning	gym	tv
accommodates	1.000000	0.811684	0.791003	0.124014	0.109612	0.253955
bedrooms	0.811684	1.000000	0.852407	0.084183	-0.034519	0.109999
beds	0.791003	0.852407	1.000000	0.105523	-0.030464	0.086498
air_conditioning	0.124014	0.084183	0.105523	1.000000	0.006799	0.136304
gym	0.109612	-0.034519	-0.030464	0.006799	1.000000	0.247503
tv	0.253955	0.109999	0.086498	0.136304	0.247503	1.000000
bbq	0.062040	-0.041420	-0.014476	0.057425	0.472893	0.119941
bed_linen	0.274648	0.065092	0.055282	0.228350	0.365264	0.439707
coffee_machine	0.323299	0.093667	0.080691	0.209567	0.418689	0.394103
cooking_basics	0.247697	0.158002	0.203304	-0.103488	-0.031880	0.117889
white_goods	0.044392	0.050702	0.030545	0.092243	0.188861	0.156572
elevator	0.132249	0.005295	-0.031894	0.223253	0.600709	0.257625
parking	-0.086699	-0.045489	0.051400	-0.011805	-0.299307	-0.252522
hot_tub_sauna_or_pool	0.113075	0.019042	0.027351	-0.209335	0.493113	0.127826

```
In [40]: boston3 = boston3.loc[boston3['price'] < 5000]
```

```
In [41]: boston3['price'].describe()
```

```
Out[41]: count      107382.000000
mean         147.484206
std          106.597849
min           0.000000
25%           50.000000
50%          149.000000
75%          212.000000
max          3999.000000
Name: price, dtype: float64
```

In [42]: `boston3.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107382 entries, 0 to 107388
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   price                                107382 non-null  int64
1   minimum_nights                       107382 non-null  int64
2   number_of_reviews                    107382 non-null  int64
3   reviews_per_month                   107382 non-null  float64
4   calculated_host_listings_count       107382 non-null  int64
5   availability_365                     107382 non-null  int64
6   accommodates                         107382 non-null  int64
7   bedrooms                             107382 non-null  float64
8   beds                                 107382 non-null  float64
9   air_conditioning                     107382 non-null  float64
10  gym                                   107382 non-null  float64
11  tv                                    107382 non-null  float64
12  bbq                                   107382 non-null  float64
13  bed_linen                             107382 non-null  float64
14  coffee_machine                       107382 non-null  float64
15  cooking_basics                       107382 non-null  float64
16  white_goods                           107382 non-null  float64
17  elevator                             107382 non-null  float64
18  parking                              107382 non-null  float64
19  hot_tub_sauna_or_pool                107382 non-null  float64
20  private_entrance                     107382 non-null  float64
21  toiletries                           107382 non-null  float64
22  workspace                             107382 non-null  float64
23  refrigerator                          107382 non-null  float64
dtypes: float64(18), int64(6)
memory usage: 20.5 MB
```

In [43]: `from sklearn.model_selection import train_test_split`
`x_train, x_test, y_train, y_test = train_test_split(boston3.drop(["price"],`

In [44]: `from ALY6040_Group3 import normalize_data`
`xscaler, x_processed = normalize_data(x_train)`

Decision Tree

```
In [45]: from sklearn.tree import DecisionTreeRegressor
from ALY6040_Group3 import train_model,score,predict

modell = decisionTree = train_model(x_processed,y_train,DecisionTreeRegressor)
start_time = time.time()
e = predict(modell,xscaler,x_test)
o = score(y_test,e)
print(type(modell).__name__,o)
print("Execution time: " + str((time.time() - start_time)) + ' ms')
a,b,c = o
xc=str((time.time() - start_time))

tab = pd.DataFrame({'Actual Values': np.array(y_test).flatten(), 'Decision
tab.set_index('Actual Values', inplace=True)
tab
```

DecisionTreeRegressor {'MAE': 11.325720752933316, 'R2': 0.8473190738130825, 'RMSE': 41.55910163033422}
Execution time: 0.020350217819213867 ms

Out[45]:

Decision Tree	
Actual Values	
34	33.969697
203	201.531915
240	240.000000
128	128.000000
131	131.000000
341	341.000000
40	42.285714
247	247.000000
52	52.000000
34	35.338462
299	299.000000
196	196.000000
177	214.333333
120	120.000000
181	181.000000
205	179.800000
194	187.666667
132	132.000000
277	277.000000
60	60.000000

In [46]: *### Predicted Values vs Actual Values in Decision Tree*

```
plt.figure(figsize=(15,8))
ax1=sb.distplot(y_test,hist=False,color='blue',label='real value')
ax2=sb.distplot(e,hist=False,color='red',label='predicted value')
plt.legend()
```


/Users/miravparekh/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

/Users/miravparekh/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

Out[46]: <matplotlib.legend.Legend at 0x7f8889bc4c70>




Feature Importance

In [47]: *### Feature importance plot of Decision Tree*

```
def plot_feature_importance(importance, names, model_type):
    feature_importance = np.array(importance)
    feature_names = np.array(names)
    data = {'feature_names': feature_names, 'feature_importance': feature_importance}
    fi_df = pd.DataFrame(data)
    fi_df.sort_values(by=['feature_importance'], ascending=False, inplace=True)
    plt.figure(figsize=(10, 8))
    sb.barplot(x=fi_df['feature_importance'], y=fi_df['feature_names'])
    plt.title(model_type)
    plt.xlabel('FEATURE IMPORTANCE')
    plt.ylabel('FEATURE NAMES')

plot_feature_importance(model1.feature_importances_, x_train.columns, 'Decision Tree')
```



Cross Validation

```
In [48]: from sklearn.model_selection import cross_val_score
```

```
lm3 = DecisionTreeRegressor()  
scores = cross_val_score(lm3, x_train, y_train, scoring='r2', cv=7)  
print('scores:', scores)  
print('Mean Score', np.mean(scores))
```

```
scores: [0.87907627 0.83702184 0.82045976 0.86588401 0.89675677 0.7702325  
4  
0.88322284]  
Mean Score 0.8503791470722932
```

Random Forest Regressor

```
In [49]: from sklearn.ensemble import RandomForestRegressor

model2 = RandomForestRegressor = train_model(x_processed,y_train,RandomFore
start_time = time.time()
u = predict(model2,xscaler,x_test)
r = score(y_test,u)
print(type(model2).__name__,r)
print("Execution time: " + str((time.time() - start_time)) + ' ms')

xcl=str((time.time() - start_time))

tab['Random Forest Regressor'] = np.array(u[:20])
tab
```

RandomForestRegressor {'MAE': 11.190295587176706, 'R2': 0.9093543087804413, 'RMSE': 32.021917965836806}
 Execution time: 0.6434817314147949 ms

Out[49]:

	Decision Tree	Random Forest Regressor
Actual Values		
34	33.969697	33.925001
203	201.531915	201.523025
240	240.000000	240.000000
128	128.000000	128.000000
131	131.000000	131.180000
341	341.000000	341.000000
40	42.285714	42.209118
247	247.000000	247.000000
52	52.000000	52.000000
34	35.338462	35.395881
299	299.000000	299.000000
196	196.000000	196.000000
177	214.333333	214.460994
120	120.000000	120.000000
181	181.000000	181.050000
205	179.800000	176.761500
194	187.666667	186.478005
132	132.000000	132.000000
277	277.000000	277.000000
60	60.000000	59.980000

In [50]: *### Predicted Values vs Actual Values in Random Forest Regressor*

```
plt.figure(figsize=(15,8))
ax1=sb.distplot(y_test,hist=False,color='blue',label='real value')
ax2=sb.distplot(u,hist=False,color='red',label='predicted value')
plt.legend()
```


/Users/miravparekh/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

/Users/miravparekh/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

Out[50]: <matplotlib.legend.Legend at 0x7f893d41f2b0>

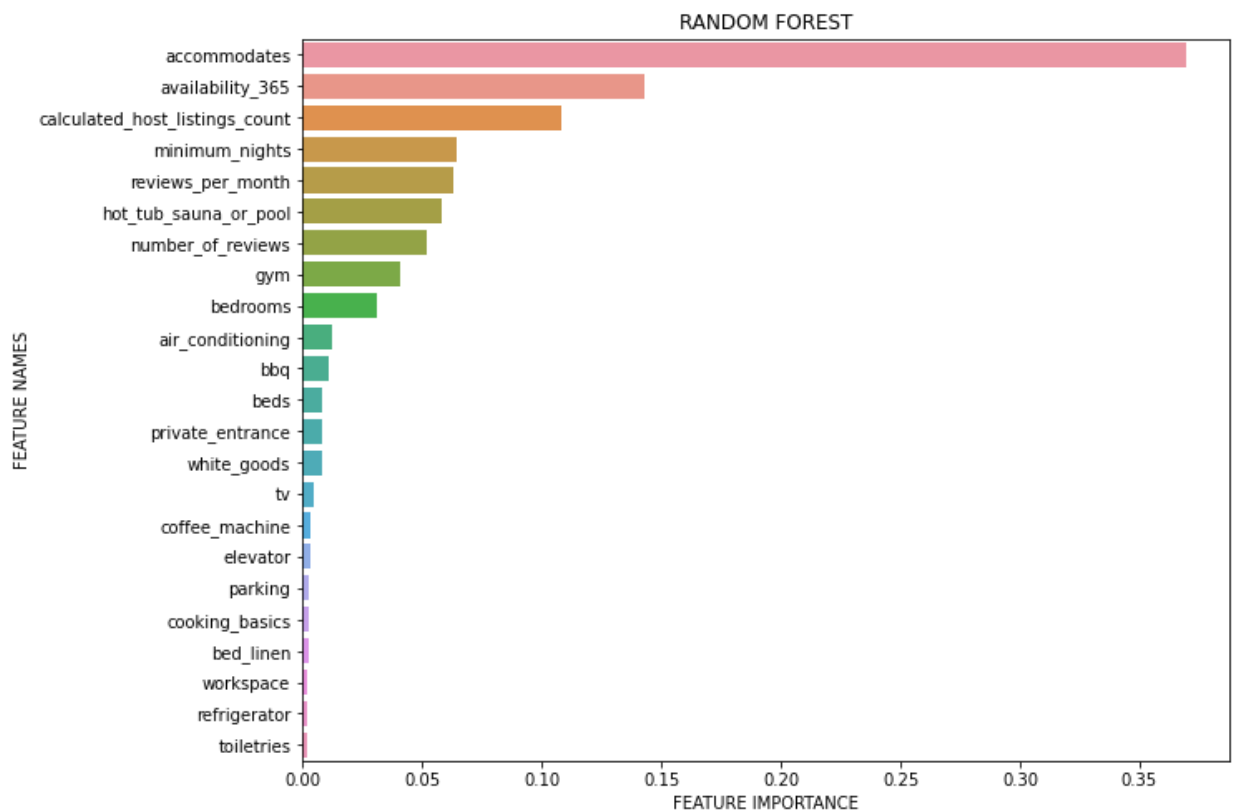


Feature Importance

```
In [51]: def plot_feature_importance(importance, names, model_type):

    feature_importance = np.array(importance)
    feature_names = np.array(names)
    data = {'feature_names': feature_names, 'feature_importance': feature_importance}
    fi_df = pd.DataFrame(data)
    fi_df.sort_values(by=['feature_importance'], ascending=False, inplace=True)
    plt.figure(figsize=(10, 8))
    sb.barplot(x=fi_df['feature_importance'], y=fi_df['feature_names'])
    plt.title(model_type)
    plt.xlabel('FEATURE IMPORTANCE')
    plt.ylabel('FEATURE NAMES')

plot_feature_importance(model2.feature_importances_, x_train.columns, 'RANDOM
```



Hyperparameter Tuning

```
In [52]: from sklearn.ensemble import RandomForestRegressor
Rf = RandomForestRegressor(random_state=42)
Rf.fit(x_train, y_train)
```

```
Out[52]: RandomForestRegressor(random_state=42)
```

```
In [53]: from sklearn.model_selection import GridSearchCV
param_grid = {
    'bootstrap': [True],
    'max_depth': [80, 85, 90],
    'max_features': ['auto'],
    'min_samples_leaf': [1],
    'min_samples_split': [2, 4],
    'n_estimators': [780, 800, 820]
}
```

```
In [54]: grid_search = GridSearchCV(estimator = Rf,
                                   param_grid = param_grid,
                                   cv = 2, n_jobs = -1, verbose = 2,
                                   scoring = 'accuracy')
```

```
In [55]: grid_search.fit(x_train, y_train)
grid_search.best_params_
```

Fitting 2 folds for each of 18 candidates, totalling 36 fits

```
-----
--
KeyboardInterrupt                                Traceback (most recent call las
t)
<ipython-input-55-a52e1defc50d> in <module>
----> 1 grid_search.fit(x_train, y_train)
      2 grid_search.best_params_

~/opt/anaconda3/lib/python3.8/site-packages/sklearn/utils/validation.py i
n inner_f(*args, **kwargs)
      61         extra_args = len(args) - len(all_args)
      62         if extra_args <= 0:
--> 63             return f(*args, **kwargs)
      64
      65         # extra_args > 0

~/opt/anaconda3/lib/python3.8/site-packages/sklearn/model_selection/_sear
ch.py in fit(self, X, y, groups, **fit_params)
      839         return results
      840
--> 841         self._run_search(evaluate_candidates)
      842
      843         # multimetric is determined here because in the case
of a callable

~/opt/anaconda3/lib/python3.8/site-packages/sklearn/model_selection/_sear
ch.py in _run_search(self, evaluate_candidates)
     1286     def _run_search(self, evaluate_candidates):
     1287         """Search all candidates in param_grid"""
-> 1288         evaluate_candidates(ParameterGrid(self.param_grid))
     1289
     1290

~/opt/anaconda3/lib/python3.8/site-packages/sklearn/model_selection/_sear
ch.py in evaluate_candidates(candidate_params, cv, more_results)
      793         n_splits, n_candidates, n_candidate
s * n_splits))
      794
--> 795         out = parallel(delayed(_fit_and_score)(clone(base
_estimator),
      796                                             X, y,
      797                                             train=train,
n, test=test,

~/opt/anaconda3/lib/python3.8/site-packages/joblib/parallel.py in __call_
_(self, iterable)
     1052
     1053         with self._backend.retrieval_context():
-> 1054             self.retrieve()
     1055         # Make sure that we get a last message telling us we
are done
     1056         elapsed_time = time.time() - self._start_time
```



```

~/opt/anaconda3/lib/python3.8/site-packages/joblib/parallel.py in retrieve(self)
    931         try:
    932             if getattr(self._backend, 'supports_timeout', False):
--> 933                 self._output.extend(job.get(timeout=self.timeout))
    934             else:
    935                 self._output.extend(job.get())

~/opt/anaconda3/lib/python3.8/site-packages/joblib/_parallel_backends.py
in wrap_future_result(future, timeout)
    540     AsyncResults.get from multiprocessing."""
    541     try:
--> 542         return future.result(timeout=timeout)
    543     except CfTimeoutError as e:
    544         raise TimeoutError from e

~/opt/anaconda3/lib/python3.8/concurrent/futures/_base.py in result(self, timeout)
    432         return self.__get_result()
    433
--> 434         self._condition.wait(timeout)
    435
    436         if self._state in [CANCELLED, CANCELLED_AND_NOTIFIED]
:

~/opt/anaconda3/lib/python3.8/threading.py in wait(self, timeout)
    300     try: # restore state no matter what (e.g., KeyboardInterrupt)
    301         if timeout is None:
--> 302             waiter.acquire()
    303             gotit = True
    304         else:

KeyboardInterrupt:

```

```

In [ ]: from sklearn.ensemble import RandomForestRegressor

Rf = RandomForestRegressor(random_state=42, bootstrap= True,
    max_depth= 80,
    max_features= 'auto',
    min_samples_leaf= 1,
    min_samples_split= 2,
    n_estimators= 780)

Rf.fit(x_train, y_train)

```

```

In [ ]: ss= Rf.predict(x_test)
    ssg = score(y_test,ss)
    print(type(Rf).__name__,ssg)

```

Cross Validation

```
In [ ]: from sklearn.model_selection import cross_val_score
        from sklearn.ensemble import RandomForestRegressor

        lm2 = RandomForestRegressor()
        scores = cross_val_score(lm2, x_train, y_train, scoring='r2', cv=7)
        print('scores:', scores)
        print('Mean Score', np.mean(scores))
```

XGBoost

```
In [56]: import xgboost as xgb

model3 = xgb = train_model(x_processed,y_train,xgb.XGBRegressor())
start_time = time.time()
h = predict(model3,xscaler,x_test)
j = score(y_test,h)
print(type(model3).__name__,j)
print("Execution time: " + str((time.time() - start_time)) + ' ms')

xc2=str((time.time() - start_time))

tab['XGBoost'] = np.array(h[:20])
tab
```

/Users/miravparekh/opt/anaconda3/lib/python3.8/site-packages/xgboost/comp
at.py:31: FutureWarning:

pandas.Int64Index is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

XGBRegressor {'MAE': 15.525294087041559, 'R2': 0.9120283150987354, 'RMSE': 31.546066398182603}
Execution time: 0.055891990661621094 ms

Out[56]:

	Decision Tree	Random Forest Regressor	XGBoost
Actual Values			
34	33.969697	33.925001	35.097530
203	201.531915	201.523025	196.505692
240	240.000000	240.000000	230.773285
128	128.000000	128.000000	151.425720
131	131.000000	131.180000	131.134521
341	341.000000	341.000000	328.379211
40	42.285714	42.209118	45.962440
247	247.000000	247.000000	233.773773
52	52.000000	52.000000	48.915230
34	35.338462	35.395881	38.652424
299	299.000000	299.000000	310.659882
196	196.000000	196.000000	190.672577
177	214.333333	214.460994	214.501129
120	120.000000	120.000000	130.386276
181	181.000000	181.050000	185.921860
205	179.800000	176.761500	178.372894
194	187.666667	186.478005	193.043060
132	132.000000	132.000000	156.580856

	Decision Tree	Random Forest Regressor	XGBoost
Actual Values			
277	277.000000	277.000000	252.791519
60	60.000000	59.980000	67.213402


```
In [ ]: plt.figure(figsize=(15,8))
ax1=sb.distplot(y_test,hist=False,color='blue',label='real value')
ax2=sb.distplot(u,hist=False,color='red',label='predicted value')
plt.legend()
```

Feature Importance

```
In [57]: def plot_feature_importance(importance,names,model_type):

    feature_importance = np.array(importance)
    feature_names = np.array(names)
    data={'feature_names':feature_names,'feature_importance':feature_importance}
    fi_df = pd.DataFrame(data)
    fi_df.sort_values(by=['feature_importance'], ascending=False,inplace=True)
    plt.figure(figsize=(10,8))
    sb.barplot(x=fi_df['feature_importance'], y=fi_df['feature_names'])
    plt.title(model_type)
    plt.xlabel('FEATURE IMPORTANCE')
    plt.ylabel('FEATURE NAMES')

    plot_feature_importance(model3.feature_importances_,x_train.columns,'XGBoost')
```



Hyperparameter Tuning

```
In [ ]: import xgboost as xgb
        xg = xgb.XGBRegressor(random_state=42)
        xg.fit(x_train, y_train)
```

```
In [ ]: from sklearn.model_selection import GridSearchCV
        param_grid = {'learning_rate': [0.1, 0.05],
                       'max_depth': [5, 7, 9],
                       'n_estimators': [100, 500, 900]}
```

```
In [ ]: gs = GridSearchCV(estimator = xg,
                          param_grid = param_grid,
                          cv = 2, n_jobs = -1, verbose = 2,
                          scoring = 'accuracy')
```

```
In [ ]: gs.fit(x_train, y_train)
        gs.best_params_
```

```
In [ ]: from xgboost import XGBRegressor
        xg = XGBRegressor(random_state = 42, learning_rate = 0.1, max_depth = 5, n_
        xg.fit(x_train,y_train)
```

```
In [ ]: xj= xg.predict(x_train)
        dsx = score(y_train,xj)
        print(type(xg).__name__,dsx)
```

Cross Validation

```
In [ ]: from sklearn.model_selection import cross_val_score
        import xgboost as xgb
        lm1 = xgb.XGBRegressor()
        scores = cross_val_score(lm1, x_train, y_train, scoring='r2', cv=7)
        print('scores:',scores)
        print('Mean Score',np.mean(scores))
```

Linear Regression

```
In [58]: from sklearn.linear_model import LinearRegression

model4 = linearRegression = train_model(x_processed,y_train,LinearRegression)
start_time = time.time()
w = predict(model4,xscaler,x_test)
p = score(y_test,w)
print(type(model4).__name__,p)
print("Execution time: " + str((time.time() - start_time)) + ' ms')

xc3=str((time.time() - start_time))

tab['Linear Regression'] = np.array(w[:20])
tab
```

LinearRegression {'MAE': 41.337713933599446, 'R2': 0.4810204575904966, 'RMSE': 76.6211873143184}
 Execution time: 0.013142108917236328 ms

Out[58]:

	Decision Tree	Random Forest Regressor	XGBoost	Linear Regression
Actual Values				
34	33.969697	33.925001	35.097530	59.999143
203	201.531915	201.523025	196.505692	207.851565
240	240.000000	240.000000	230.773285	203.099159
128	128.000000	128.000000	151.425720	193.862797
131	131.000000	131.180000	131.134521	154.622943
341	341.000000	341.000000	328.379211	206.177947
40	42.285714	42.209118	45.962440	47.359935
247	247.000000	247.000000	233.773773	210.418266
52	52.000000	52.000000	48.915230	70.707581
34	35.338462	35.395881	38.652424	103.220043
299	299.000000	299.000000	310.659882	255.795042
196	196.000000	196.000000	190.672577	206.727563
177	214.333333	214.460994	214.501129	210.148438
120	120.000000	120.000000	130.386276	187.213011
181	181.000000	181.050000	185.921860	199.358890
205	179.800000	176.761500	178.372894	205.837079
194	187.666667	186.478005	193.043060	203.354411
132	132.000000	132.000000	156.580856	198.270749
277	277.000000	277.000000	252.791519	172.610189
60	60.000000	59.980000	67.213402	143.253494

Cross Validation

```
In [ ]: from sklearn.model_selection import cross_val_score

lm = LinearRegression()
scores = cross_val_score(lm, x_train, y_train, scoring='r2', cv=7)
print('scores:', scores)
print('Mean Score', np.mean(scores))
```

```
In [ ]: plt.figure(figsize=(15,8))
ax1=sb.distplot(y_test,hist=False,color='blue',label='real value')
ax2=sb.distplot(u,hist=False,color='red',label='predicted value')
plt.legend()
```

Feature Importance

```
In [ ]: importance = model4.coef_
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))
    plt.xlabel('FEATURE IMPORTANCE')
    plt.ylabel('FEATURE NAMES')
plot_feature_importance([x for x in range(len(importance))],importance,'Lin
```

Regularization

```
In [ ]: # Creating a linear regression model with Elastic net
re = ElasticNet(alpha=1.0,l1_ratio=0.5)
re.fit(boston3.drop(["price"],axis=1),boston3.price)
y_pred=re.predict(x_test)

print("RMSE: %.3f" % mean_squared_error(y_test, y_pred))
print("R2: %.3f" % r2_score(y_test, y_pred))
```

```
In [ ]: import statsmodels.formula.api as smf

#fit regression model
fit = smf.ols('price ~ host_id+minimum_nights+number_of_reviews+reviews_per

#view model summary
print(fit.summary())

from statsmodels.compat import lzip
import statsmodels.stats.api as sms

#perform Bresuch-Pagan test
names = ['Lagrange multiplier statistic', 'p-value',
         'f-value', 'f p-value']
test = sms.het_breuschpagan(fit.resid, fit.model.exog)

lzip(names, test)
```

```
In [ ]: from tabulate import tabulate

tab1 = [
    ["Decision Tree", xc],
    ["Random Forest Regressor", xc1],
    ["XGBoost", xc2],
    ["Linear Regression", xc3]
]

head = ["Model", "Execution Time"]
print(tabulate(tab1, headers=head, tablefmt="grid"))
```



```
In [ ]: from tabulate import tabulate

xz=o[ 'RMSE' ]
xz1=r[ 'RMSE' ]
xz2=dsx[ 'RMSE' ]
xz3=p[ 'RMSE' ]

xxz=o[ 'MAE' ]
xxz1=r[ 'MAE' ]
xxz2=dsx[ 'MAE' ]
xxz3=p[ 'MAE' ]

xxxz=o[ 'R2' ]
xxxz1=r[ 'R2' ]
xxxz2=dsx[ 'R2' ]
xxxz3=p[ 'R2' ]

tab2 = [
    ["Decision Tree",xxz,xxxz,xz],
    ["Random Forest Regressor",xxz1,xxxz1,xz1],
    ["XGBoost",xxz2,xxxz2,xz2],
    ["Linear Regression",xxz3,xxxz3,xz3]
]

head = [ "Model", "MAE", "R2", "RMSE" ]
print(tabulate(tab2, headers=head, tablefmt="grid"))
```

```
In [ ]:
```